

`ranges::to`: A function to convert any range to a container

Document #: P1206R1  
Date: 2019-01-20  
Project: Programming Language C++  
Audience: LEWG  
Reply-to: Corentin Jabot <[corentin.jabot@gmail.com](mailto:corentin.jabot@gmail.com)>  
Eric Niebler <[eric.niebler@gmail.com](mailto:eric.niebler@gmail.com)>  
Casey Carter <[Casey@carter.net](mailto:Casey@carter.net)>  
Christopher Di Bella <[cjdb.ns@gmail.com](mailto:cjdb.ns@gmail.com)>

## 1 Abstract

We propose a function to copy or materialize any range (containers and views alike) to a container.

## 2 Revisions

### Revision 1

- Split out the proposed constructors for string view and span into separate papers ([[P1391](#)] and [[P1394](#)] respectively)
- Use a function based approach rather than adding a constructor to standard containers, as it proved unworkable.

### 3 Quick Overview

We propose all the following syntaxes to be valid constructs

```
std::list<int> l;
std::map<int, int> m;

// copy a list to a vector of the same type
Same<std::vector<int>> auto a = ranges::to<std::vector<int>>(l);
//Specify an allocator
Same<std::vector<int, Alloc>> auto b = ranges::to<std::vector<int, Alloc>(l, alloc);
// copy a list to a vector of the same type, deducing value_type
Same<std::vector<int>> auto c = ranges::to<std::vector>(l);
// copy to a container of types ConvertibleTo
Same<std::vector<long>> auto d = ranges::to<std::vector<long>>(l);

//Supports converting associative container to sequence containers
Same<std::vector<std::pair<const int, int>>>
    auto f = ranges::to<vector<std::pair<const int, int>>>(m);

//Removing the const from the key by default
Same<std::vector<std::pair<int, int>>> auto e = ranges::to<vector>(m);

//Supports converting sequence containers to associative ones
Same<std::map<int, int>> auto g = f | ranges::to<map>();

//Pipe syntax
Same<std::vector<int>> auto g = 1 | ranges::view::take(42) | ranges::to<std::vector>();

//Pipe syntax with allocator
auto h = 1 | ranges::view::take(42) | ranges::to<std::vector>(alloc);

//The pipe syntax also support specifying the type and conversions
auto i = 1 | ranges::view::take(42) | ranges::to<std::vector<long>>();

//Pathentheses are optional for template
Same<std::vector<int>> auto j = 1 | ranges::view::take(42) | ranges::to<std::vector>;
//and types
auto k = 1 | ranges::view::take(42) | ranges::to<std::vector<long>>;
```

## 4 Tony tables

Before	After
<pre>std::list&lt;int&gt; lst = /*...*/; std::vector&lt;int&gt; vec     {std::begin(lst), std::end(lst)};</pre>	<pre>std::vector&lt;int&gt; = lst   ranges::to&lt;std::vector&gt;;</pre>
<pre>auto view = ranges::iota(42); vector &lt;     iter_value_t&lt;         iterator_t&lt;decltype(view)&gt;     &gt; &gt; vec; if constexpr(SizedRanged&lt;decltype(view)&gt;) {     vec.reserve(ranges::size(view)); } ranges::copy(view, std::back_inserter(vec));</pre>	<pre>auto vec = ranges::iota(42)       ranges::to&lt;std::vector&gt;;</pre>
<pre>std::map&lt;int, widget&gt; map = get_widgets_map(); std::vector&lt;     typename decltype(map)::value_type &gt; vec; vec.reserve(map.size()); ranges::move(map, std::back_inserter(vec));</pre>	<pre>auto vec = get_widgets_map()       ranges::to&lt;vector&gt;;</pre>

## 5 Motivation

Most containers of the standard library provide a constructors taking a pair of iterators.

```
std::list<int> lst;
std::vector<int> vec{std::begin(lst), std::end(lst)};
//equivalent too
std::vector<int> vec;
std::copy(it, end, std::back_inserter(vec));
```

While, this feature is very useful, as converting from one container type to another is a frequent use-case, it can be greatly improved by taking full advantage of the notions and tools offered by ranges.

Indeed, given all containers are ranges (ie: an iterator-sentinel pair) the above example can be rewritten, without semantic as:

```
std::list<int> lst;
std::vector<int> vec = lst | ranges::to<std::vector>;
```

The above example is a common pattern as it is frequently preferable to copy the content of a `std::list` to a `std::vector` before feeding it an algorithm and then copying it back to a `std::vector`.

As all containers and views are ranges, it is logical they can themselves be easily built out of ranges.

## 5.1 View Materialization

The main motivation for this proposal is what is colloquially called *view materialization*. A view can generate its elements lazily (upon increment or decrement), such as the value at a given position of the sequence iterated over only exist transiently in memory if an iterator is pointing to that position. (Note: while all lazy ranges are views, not all views are lazy).

*View materialization* consists in committing all the elements of such view in memory by putting them into a container.

The following code iterates over the numbers 0 to 1023 but only one number actually exists in memory at any given time.

```
std::iota_view v{0, 1024};
for (auto i : v) {
    std::cout << i << ' ';
}
```

While this offers great performance and reduced memory footprint, it is often necessary to put the result of the transformation operated by the view into memory. The facilities provided by [\[P0896R3\]](#) allow to do that in the following way:

```
std::iota_view v{0, 1024};
std::vector<int> materialized;
std::copy(v, std::back_inserter(materialized));
```

This proposal allows rewriting the above snippet as:

```
auto materialized = std::iota_view{0, 1024} | std::ranges::to<std::vector>();
```

Perhaps the most important aspect of view materialization is that it allows simple code such as:

```
namespace std {
    split_view<std::string_view> split(std::string_view);
}
auto res = std::split("Splitting strings made easy")
    | std::ranges::to<std::vector>;
```

Indeed, a function such as `split` is notoriously hard to standardize ([\[P0540\]](#), [\[N3593\]](#)), because without lazy views and `std::string_view`, it has to allocate or expose an expert-friendly interface. The view materialization pattern further let the *caller* choose the best container and allocation strategy for their use case (or to never materialize the view should it not be necessary). And while it would not make sense for a standard-library function to split a string into a vector it would allocate, it's totally reasonable for most applications to do so.

This paper does not propose to standardize such `split` function - a `split_view` exist in [\[P0896R3\]](#), however, view materialization is something the SG-16 working group is interested in. Indeed, they

have considered APIs that could rely heavily on this idiom, as it has proven a natural way to handle the numerous ways to iterate over Unicode text. Similar ideas have been presented in [P1004].

```
auto sentences =
    text(blob)
    normalize<text::nfc> |
    graphemes_view |
    split<sentences> | ranges::to<std::vector<std::u8string>>;
```

## 6 Constructing views from ranges

Constructing standard views (`string_view` and `span`) from ranges is addressed in separate papers as the design space and the requirements are different:

- `string_view` : [P1391]
- `span` : [P1394]
- Work is being done to allow Ranges's iterators to be move only

As views are not containers, they are not constructible from `ranges::to`

## 7 Alternative designs

While we believe the range constructor based approach is the cleanest way to solve this problem, LEWG was interested in alternative design based on free functions

### 7.1 Range constructors

The original version of that paper proposed to add range constructors to all constructors. This proved to be unworkable because of `std::initializer_list`:

```
std::vector<int> foo = ....;
std::vector a{foo}; //constructs a std::vector<std::vector<int>>
std::vector b(foo); //would construct a std::vector<int>
```

## 8 Existing practices

### 8.1 Range V3

This proposal is based on the `to` (previously `(to_)` function offered by ranges v3.

```

auto vec = view::ints
| view::transform([](int i) {
    return i + 42;
})
| view::take(10)
| to<std::vector>;

```

## 8.2 Abseil

Abseil offer converting constructors with each of their view. As per their documentation:

One of the more useful features of the `StrSplit()` API is its ability to adapt its result set to the desired return type. `StrSplit()` returned collections may contain `std::string`, `absl::string_view`, or any object that can be explicitly created from an `absl::string_view`. This pattern works for all standard STL containers including `std::vector`, `std::list`, `std::deque`, `std::set`, `std::multiset`, `std::map`, and `std::multimap`, and even `std::pair`, which is not actually a container.

Because they can not modify existing containers, view materialization in Abseil is done by the mean of a conversion operator:

```

template<Container C>
operator C();

```

However, because it stands to reason to expect that there are many more views than containers and because conversions between containers are also useful, it is a more general solution to provide a solution that is not coupled with each individual view.

## 8.3 Previous work

[\[N3686\]](#) explores similar solutions and was discussed by LEWG long before the Ranges TS.

## 9 Future work

Eric Niebler suggested (and implemented in ranges v3) making the parenthesis `range::to<...>()` in optional, such that

```

auto v = view | range::to<vector<int>>;

```

Would be valid syntax. However, this is currently only be feasible for types `template (range::to<vector<int>>;)` and not `template template (range::to<vector>;)` which would require a core language tweak to make possible.

## 10 Proposed wording

We do not provide wording at this time, but this is what the interface would look like conceptually.

```
namespace ranges {

    //1
    template <Container C, Range R, typename...Arg>
    constexpr auto to(const R & r, Args...&) -> C;

    //2
    template <template <typename...> typename C,
              Range R, typename T = range_value_t<R>, typename... Args>
    constexpr auto to(const R & r, Args...&) -> C<T, Args...>;

    //3
    template <Container C, typename...Args>
    constexpr auto to(Args...&&) -> implementation-defined;

    //4
    template <template <typename...> typename C>
    constexpr auto to(Args...&&) -> implementation-defined;

    //5
    template <Range R>
    constexpr auto operator|(const R && r, implementation-defined{});
}
```

Functions 3 and 4 return an implementation defined object that can be passed to the pipe operator(5), and provide an implementation defined way of creating a container of the appropriate type.

`range::to` forwards the range to the container if it is constructible from it. If not, it tries to construct the container from the `begin(range)/end(range)` iterators pair. Otherwise, it falls-back to `ranges::copy`. This ensure the most efficient strategy is selected to perform the actual copy.

## 11 Implementation Experience

Implementations of this proposal are available on in the 1.0 branch of [\[RangeV3\]](#) and on [\[Compiler Explorer\]](#). To make sure the parentheses are optional (`v | ranges::to<vector>;`) our implementations use a default constructed tag which dispatch through a function pointer. However, this have no run-time cost and doesn't suffer from the same issues LEWG had about `std::in_place_tag` because no actual indirection takes place. We believe being able to omit the parenthesis is necessary so `ranges::to` remains consistent with the syntax of views adaptors, and is otherwise a nice quality of life improvement in a facility which we expect to be used frequently.

## 12 Acknowledgements

We would like to thank the people who gave feedback on this paper, notably Arthur O’Dwyer, Barry Revzin and Tristan Brindle.

## 13 References

- [Compiler Explorer] <https://godbolt.org/z/m1I9NE>
- [RangeV3] Eric Niebler [https://github.com/ericniebler/range-v3/blob/v1.0-beta/include/range/v3/to\\_container.hpp](https://github.com/ericniebler/range-v3/blob/v1.0-beta/include/range/v3/to_container.hpp)
- [P1391] Corentin Jabot *Range constructor for std::string\_view*  
<https://wg21.link/P1391>
- [P1394] Corentin Jabot *Range constructor for std::span*  
<https://wg21.link/P1394>
- [P0896R3] Eric Niebler, Casey Carter, Christopher Di Bella *The One Range Ts Proposal*  
<https://wg21.link/P0896>
- [P1004] Louis Dionne *Making std::vector constexpr*  
<https://wg21.link/P1004>
- [P1004] Tom Honermann *Text\_view: A C++ concepts and range based character encoding and code point enumeration library*  
<https://wg21.link/P0244>
- [P0540] Laurent Navarro *A Proposal to Add split/join of string/string\_view to the Standard Library*  
<https://wg21.link/P0540>
- [N3593] Greg Miller *std::split(): An algorithm for splitting strings*  
<https://wg21.link/N3593>
- [P1035] Christopher Di Bella *Input range adaptors*  
<https://wg21.link/P1035>
- [Abseil] <https://abseil.io/docs/cpp/guides/strings>
- [N3686] Jeffrey Yasskin *[Ranges] Traversable arguments for container constructors and methods*  
<https://wg21.link/n3686>
- [P1072R1] Chris Kennelly, Mark Zeren *Vector as allocation transfer device* <https://wg21.link/P1072>
- [1] [P0504R0] Jonathan Wakely *Revisiting in-place tag types for any/optional/variant* <https://wg21.link/P0504R0>